

Bitwise Operators

C also provides *bitwise* operators which work with bits.

C bitwise operators: `& | ^ ~ << >>`

Bitwise AND

The & operator

- takes two values (1,2,4,8 bytes), treats as sequence of bits
- performs logical AND on each corresponding pair of bits
- result contains same number of bits as inputs

Example:

00100111	AND		0	1
& 11100011	----		-----	
-----			0	0 0
00100011			1	0 1

Used for e.g. checking whether a bit is set

Checking for Odd Numbers

The obvious way to check for odd numbers in C

```
int isOdd(int n) {  
    return n % 2 == 1;  
}
```

We can use & to achieve the same thing:

```
int isOdd(int n) {  
    return n & 1;  
}
```

Bitwise OR

The `|` operator

- takes two values (1,2,4,8 bytes), treats as sequence of bits
- performs logical OR on each corresponding pair of bits
- result contains same number of bits as inputs

Example:

00100111	OR		0	1
11100011	----		-----	
-----	0		0	1
11100111	1		1	1

Used for e.g. ensuring that a bit is set

Bitwise NEG

The ~ operator

- takes a single value (1,2,4,8 bytes), treats as sequence of bits
- performs logical negation of each bit
- result contains same number of bits as input

Example:

~ 00100111	NEG 0 1
-----	---- -----
11011000	1 0

Used for e.g. creating useful bit patterns

Bitwise Operations in C

- everything is ultimately a string of bits
- e.g. `unsigned char` = 8-bit value
- e.g. literal bit-string `0b01110001`
- e.g. literal hexadecimal `0x71`
- `&` = bitwise AND
- `|` = bitwise OR
- `~` = bitwise NEG

Bitwise XOR

The ^ operator

- takes two values (1,2,4,8 bytes), treats as sequence of bits
- performs logical XOR on each corresponding pair of bits
- result contains same number of bits as inputs

Example:

00100111	XOR		0	1
^ 11100011	----		-----	
-----	0		0	1
11000100	1		1	0

Used in e.g. generating hashes, graphic operation, cryptography

Left Shift

The << operator

- takes a single value (1,2,4,8 bytes), treats as sequence of bits
- and a small positive integer x
- moves (shifts) each bit x positions to the left
- left-end bit vanishes; right-end bit replaced by zero
- result contains same number of bits as input

Example:

00100111 << 2	00100111 << 8
-----	-----
10011100	00000000

Right Shift

The >> operator

- takes a single value (1,2,4,8 bytes), treats as sequence of bits
- and a small positive integer x
- moves (shifts) each bit x positions to the right
- right-end bit vanishes; left-end bit replaced by zero**
- result contains same number of bits as input

Example:

00100111 >> 2	00100111 >> 8
-----	-----
00001001	00000000

- shifts involving negative values are not portable (implementation defined)
- common source of bugs in COMP1521 and elsewhere
- always use unsigned values/variables to be safe/portable.

bitwise.c: showing results of bitwise operation

```
$ gcc bitwise.c print_bits.c -o bitwise
$ ./bitwise
Enter a: 23032
Enter b: 12345
Enter c: 3
      a = 0101100111111000 = 0x59f8 = 23032
      b = 0011000000111001 = 0x3039 = 12345
     ~a = 1010011000000111 = 0xa607 = 42503
a & b = 0001000000111000 = 0x1038 = 4152
a | b = 0111100111111001 = 0x79f9 = 31225
a ^ b = 0110100111000001 = 0x69c1 = 27073
a >> c = 0000101100111111 = 0x0b3f = 2879
a << c = 1100111111000000 = 0xcfc0 = 53184
```

source code for bitwise.c

source code for print_bits.c source code for print_bits.h

bitwise.c: code

```
uint16_t a = 0;
printf("Enter a: ");
scanf("%hd", &a);
uint16_t b = 0;
printf("Enter b: ");
scanf("%hd", &b);
printf("Enter c: ");
int c = 0;
scanf("%d", &c);
print_bits_hex("    a = ", a);
print_bits_hex("    b = ", b);
print_bits_hex("   ~a = ", ~a);
print_bits_hex(" a & b = ", a & b);
print_bits_hex(" a | b = ", a | b);
print_bits_hex(" a ^ b = ", a ^ b);
print_bits_hex("a >> c = ", a >> c);
print_bits_hex("a << c = ", a << c);
```

shift_as_multiply.c: using shift to multiply by 2^n

```
$ gcc shift_as_multiply.c print_bits.c -o shift_as_multiply
$ ./shift_as_multiply 4
2 to the power of 4 is 16
In binary it is: 000000000000000000000000000010000
$ ./shift_as_multiply 20
2 to the power of 20 is 1048576
In binary it is: 0000000000001000000000000000000000
$ ./shift_as_multiply 31
2 to the power of 31 is 2147483648
In binary it is: 100000000000000000000000000000000
$
```

shift_as_multiply.c: using shift to multiply by 2^n

```
int n = strtol(argv[1], NULL, 0);
uint32_t power_of_two;
int n_bits = 8 * sizeof power_of_two;
if (n >= n_bits) {
    fprintf(stderr, "n is too large\n");
    return 1;
}
power_of_two = 1;
power_of_two = power_of_two << n;
printf("2 to the power of %d is %u\n", n, power_of_two);
printf("In binary it is: ");
print_bits(power_of_two, n_bits);
printf("\n");
```

source code for shift_as_multiply.c

set_low_bits.c: using << and - to set low n bits

```
$ dcc set_low_bits.c print_bits.c -o n_ones
```

```
$ ./set_low_bits 3
```

The bottom 3 bits of 7 are ones:

[illegible]

```
$ ./set_low_bits 19
```

The bottom 19 bits of 524287 are ones:

```
0000000000000011111111111111111111
```

```
$ ./set low bits 29
```

The bottom 29 bits of 536870911 are ones:

00011111111111111111111111111111

set_low_bits.c: using << and - to set low n bits

```
int n = strtol(argv[1], NULL, 0);
uint32_t mask;
int n_bits = 8 * sizeof mask;
assert(n >= 0 && n < n_bits);
mask = 1;
mask = mask << n;
mask = mask - 1;
printf("The bottom %d bits of %u are ones:\n", n, mask);
print_bits(mask, n_bits);
printf("\n");
```

source code for set_low_bits.c

set_bit_range.c: using << and - to set a range of bits

```
$ gcc set_bit_range.c print_bits.c -o set_bit_range
$ ./set_bit_range 0 7
Bits 0 to 7 of 255 are ones:
00000000000000000000000011111111
$ ./set_bit_range 8 15
Bits 8 to 15 of 65280 are ones:
000000000000000001111111100000000
$ ./set_bit_range 8 23
Bits 8 to 23 of 16776960 are ones:
000000001111111111111111100000000
$ ./set_bit_range 1 30
Bits 1 to 30 of 2147483646 are ones:
01111111111111111111111111111110
```


set_bit_range.c: using << and - to set a range of bits

```
int low_bit = strtol(argv[1], NULL, 0);
int high_bit = strtol(argv[2], NULL, 0);
uint32_t mask;
int n_bits = 8 * sizeof mask;
```

source code for set_bit_range.c

```
int mask_size = high_bit - low_bit + 1;
mask = 1;
mask = mask << mask_size;
mask = mask - 1;
mask = mask << low_bit;
printf("Bits %d to %d of %u are ones:\n", low_bit, high_bit, mask);
print_bits(mask, n_bits);
printf("\n");
```

source code for set_bit_range.c

extract_bit_range.c: extracting a range of bits

```
$ gcc extract_bit_range.c print_bits.c -o extract_bit_range
```

```
$ ./extract_bit_range 4 7 42
```

Value 42 in binary is:

00000000000000000000000000000000101010

Bits 4 to 7 of 42 are:

0010

```
$ ./extract_bit_range 10 20 123456789
```

Value 123456789 in binary is:

00000111010110111100110100010101

Bits 10 to 20 of 123456789 are:

11011110011

extract_bit_range.c: extracting a range of bits

```
int mask_size = high_bit - low_bit + 1;
mask = 1;
mask = mask << mask_size;
mask = mask - 1;
mask = mask << low_bit;
// get a value with the bits outside the range low_bit..high_bit set to 0
uint32_t extracted_bits = value & mask;
// right shift the extracted_bits so low_bit becomes bit 0
extracted_bits = extracted_bits >> low_bit;
printf("Value %u in binary is:\n", value);
print_bits(value, n_bits);
printf("\n");
printf("Bits %d to %d of %u are:\n", low_bit, high_bit, value);
print_bits(extracted_bits, mask_size);
printf("\n");
```

source code for extract_bit_range.c

print_bits.c: extracting the n-th bit of a value

```
void print_bits(uint64_t value, int how_many_bits) {  
    // print bits from most significant to least significant  
    for (int i = how_many_bits - 1; i >= 0; i--) {  
        int bit = get_nth_bit(value, i);  
        printf("%d", bit);  
    }  
}  
  
int get_nth_bit(uint64_t value, int n) {  
    // shift the bit right n bits  
    // this leaves the n-th bit as the least significant bit  
    uint64_t shifted_value = value >> n;  
    // zero all bits except the the least significant bit  
    int bit = shifted_value & 1;  
    return bit;  
}
```

print_int_in_hex.c: print an integer in hexadecimal

- write C to print an integer in hexadecimal instead of using:

```
printf("%x", n)
```

```
$ gcc print_int_in_hex.c -o print_int_in_hex
```

```
$ ./print_int_in_hex
```

```
Enter a positive int: 42
```

```
42 = 0x0000002A
```

```
$ ./print_int_in_hex
```

```
Enter a positive int: 65535
```

```
65535 = 0x0000FFFF
```

```
$ ./print_int_in_hex
```

```
Enter a positive int: 3735928559
```

```
3735928559 = 0xDEADBEEF
```

```
$
```

source code for print_int_in_hex.c

print_int_in_hex.c: main

```
int main(void) {  
    uint32_t a = 0;  
    printf("Enter a positive int: ");  
    scanf("%u", &a);  
    printf("%u = 0x", a);  
    print_hex(a);  
    printf("\n");  
    return 0;  
}
```

source code for print_int_in_hex.c

print_int_in_hex.c: print_hex - extracting digit

```
void print_hex(uint32_t n) {  
    // sizeof return number of bytes in n's representation  
    // each byte is 2 hexadecimal digits  
    int n_hex_digits = 2 * (sizeof n);  
    // print hex digits from most significant to least significant  
    for (int which_digit = n_hex_digits - 1; which_digit >= 0; which_digit--) {  
        // shift value across so hex digit we want  
        // is in bottom 4 bits  
        int bit_shift = 4 * which_digit;  
        uint32_t shifted_value = n >> bit_shift;  
        // mask off (zero) all bits but the bottom 4 bites  
        int hex_digit = shifted_value & 0xF;  
    }  
}
```

source code for print_int_in_hex.c

print_int_in_hex.c: converting digit to ASCII

```
    // hex digit will be a value 0..15
    // obtain the corresponding ASCII value
    // "0123456789ABCDEF" is a char array
    // containing the appropriate ASCII values (+ a '\0')
    int hex_digit_ascii = "0123456789ABCDEF"[hex_digit];
    putchar(hex_digit_ascii);
}
}
```

source code for print_int_in_hex.c

int_to_hex_string.c: convert int to a string of hex digits

- Write C to convert an integer to a string containing its hexadecimal digits.

Could use the C library function `snprintf` to do this.

```
$ gcc int_to_hex_string.c -o int_to_hex_string
```

```
$ ./int_to_hex_string
```

```
$ ./int_to_hex_string
```

```
Enter a positive int: 42
```

```
42 = 0x0000002A
```

```
$ ./int_to_hex_string
```

```
Enter a positive int: 65535
```

```
65535 = 0x0000FFFF
```

```
$ ./int_to_hex_string
```

```
Enter a positive int: 3735928559
```

```
3735928559 = 0xDEADBEEF
```

```
$
```

source code for `int_to_hex_string.c`

int_to_hex_string.c: main

```
int main(void) {  
    uint32_t a = 0;  
    printf("Enter a positive int: ");  
    scanf("%u", &a);  
    char *hex_string = int_to_hex_string(a);  
    // print the returned string  
    printf("%u = 0x%s\n", a, hex_string);  
    free(hex_string);  
    return 0;  
}
```

source code for int_to_hex_string.c

int_to_hex_string.c: convert int to a string of hex digits

```
char *int_to_hex_string(uint32_t n) {  
    // sizeof return number of bytes in n's representation  
    // each byte is 2 hexadecimal digits  
    int n_hex_digits = 2 * (sizeof n);  
    // allocate memory to hold the hex digits + a terminating 0  
    char *string = malloc(n_hex_digits + 1);  
    // print hex digits from most significant to least significant  
    for (int which_digit = 0; which_digit < n_hex_digits; which_digit++)  
        // shift value across so hex digit we want  
        // is in bottom 4 bits  
        int bit_shift = 4 * which_digit;  
        uint32_t shifted_value = n >> bit_shift;  
        // mask off (zero) all bits but the bottom 4 bits  
        int hex_digit = shifted_value & 0xF;
```

source code for int_to_hex_string.c

int_to_hex_string.c: convert int to a string of hex digits

```
// hex digit will be a value 0..15  
// obtain the corresponding ASCII value  
// "0123456789ABCDEF" is a char array  
// containing the appropriate ASCII values  
int hex_digit_ascii = "0123456789ABCDEF"[hex_digit];  
string[which_digit] = hex_digit_ascii;  
}  
// 0 terminate the array  
string[n_hex_digits] = 0;  
return string;  
}
```

source code for int_to_hex_string.c

hex_string_to_int.c: convert hex digit string to int

- As an exercise write C to convert an integer to a string containing its hexadecimal digits.

Could use the C library function `strtol` to do this.

```
$ gcc hex_string_to_int.c -o hex_string_to_int
$ gcc hex_string_to_int.c -o hex_string_to_int
$ ./hex_string_to_int 2A
2A hexadecimal is 42 base 10
$ ./hex_string_to_int FFFF
FFFF hexadecimal is 65535 base 10
$ ./hex_string_to_int DEADBEEF
DEADBEEF hexadecimal is 3735928559 base 10
$
```

source code for `hex_string_to_int.c`

hex_string_to_int.c: main

```
int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <hexadecimal-number>\n", argv[0]);
        return 1;
    }
    char *hex_string = argv[1];
    uint32_t u = hex_string_to_int(hex_string);
    printf("%s hexadecimal is %u base 10\n", hex_string, u);
    return 0;
}
```

source code for hex_string_to_int.c

hex_string_to_int.c: convert array of hex digits to int

```
uint32_t hex_string_to_int(char *hex_string) {  
    uint32_t value = 0;  
    for (int which_digit = 0; hex_string[which_digit] != 0; which_d  
        int ascii_hex_digit = hex_string[which_digit];  
        int digit_as_int = hex_digit_to_int(ascii_hex_digit);  
        value = value << 4;  
        value = value | digit_as_int;  
    }  
    return value;  
}
```

source code for hex_string_to_int.c

hex_string_to_int.c: convert single hex digit to int

```
int hex_digit_to_int(int ascii_digit) {  
    if (ascii_digit >= '0' && ascii_digit <= '9') {  
        // the ASCII characters '0' .. '9' are contiguous  
        // in other words they have consecutive values  
        // so subtract the ASCII value for '0' yields the correspon  
        return ascii_digit - '0';  
    }  
    if (ascii_digit >= 'A' && ascii_digit <= 'F') {  
        // for characters 'A' .. 'F' obtain the  
        // corresponding integer for a hexadecimal digit  
        return 10 + (ascii_digit - 'A');  
    }  
    fprintf(stderr, "Bad digit '%c'\n", ascii_digit);  
    exit(1);  
}
```

source code for hex_string_to_int.c

shift_bug.c: bugs to avoid

```
// int16_t is a signed type (-32768..32767)  
// below operations are undefined for a signed type  
int16_t i;  
i = -1;  
i = i >> 1; // undefined - shift of a negative value  
printf("%d\n", i);  
i = -1;  
i = i << 1; // undefined - shift of a negative value  
printf("%d\n", i);  
i = 32767;  
i = i << 1; // undefined - left shift produces a negative value  
uint64_t j;  
j = 1 << 33; // undefined - constant 1 is an int  
j = ((uint64_t)1) << 33; // ok
```

source code for shift_bug.c

xor.c: fun with xor

```
int xor_value = strtol(argv[1], NULL, 0);
if (xor_value < 0 || xor_value > 255) {
    fprintf(stderr, "Usage: %s <xor-value>\n", argv[0]);
    return 1;
}

int c;
while ((c = getchar()) != EOF) {
    //      exclusive-or
    //      ^  | 0  1
    //      ----|-----
    //      0  | 0  1
    //      1  | 1  0
    int xor_c = c ^ xor_value;
    putchar(xor_c);
}
```

source code for xor.c

xor.c: fun with xor

```
$ echo Hello Andrew|xor 42
bOFFE
kDNXO] $ echo Hello Andrew|xor 42|cat -A
bOFFE$
kDNXO] $
$ echo Hello |xor 42
bOFFE $ echo -n 'bOFFE '|xor 42
Hello
$ echo Hello|xor 123|xor 123
Hello
$
```

pokemon.c: using an int to represent a set of values

```
#define FIRE_TYPE      0x0001
#define FIGHTING_TYPE  0x0002
#define WATER_TYPE     0x0004
#define FLYING_TYPE    0x0008
#define POISON_TYPE    0x0010
#define ELECTRIC_TYPE  0x0020
#define GROUND_TYPE    0x0040
#define PSYCHIC_TYPE   0x0080
#define ROCK_TYPE      0x0100
#define ICE_TYPE       0x0200
#define BUG_TYPE       0x0400
#define DRAGON_TYPE    0x0800
#define GHOST_TYPE     0x1000
#define DARK_TYPE      0x2000
#define STEEL_TYPE     0x4000
#define FAIRY_TYPE     0x8000
```

pokemon.c: using an int to represent a set of values

- simple example of a single integer specifying a set of values
- interacting with hardware often involves this sort of code

```
uint16_t our_pokemon = BUG_TYPE | POISON_TYPE | FAIRY_TYPE;
```

```
// example code to check if a pokemon is of a type:
```

```
if (our_pokemon & POISON_TYPE) {  
    printf("Poisonous\n"); // prints  
}  
  
if (our_pokemon & GHOST_TYPE) {  
    printf("Scary\n"); // does not print  
}
```

source code for pokemon.c

pokemon.c: using an int to represent a set of values

```
// example code to add a type to a pokemon
our_pokemon |= GHOST_TYPE;
// example code to remove a type from a pokemon
our_pokemon &= ~ POISON_TYPE;

printf(" our_pokemon type (2)\n");
if (our_pokemon & POISON_TYPE) {
    printf("Poisonous\n"); // does not print
}
if (our_pokemon & GHOST_TYPE) {
    printf("Scary\n"); // prints
}
```

source code for pokemon.c

bitset.c: using an int to represent a set of values

```
$ dcc bitset.c print_bits.c -o bitset
```

```
$ ./bitset
```

Set members can be 0-63, negative number to finish

```
Enter set a: 1 2 4 8 16 32 -1
```

```
Enter set b: 5 4 3 33 -1
```

[illegible]

b = 00000000000000000000000000000000100000000000000000000000011100

$$a = \{1, 2, 4, 8, 16, 32\}$$
$$b = \{3, 4, 5, 33\}$$

```
a union b = {1,2,3,4,5,8,16,32,33}
```

a intersection b = {4}

$$\text{cardinality}(a) = 6$$

```
is_member(42, a) = 0
```

bitset.c: main

```
printf("Set members can be 0-%d, negative number to finish\n",
      MAX_SET_MEMBER);
set a = set_read("Enter set a: ");
set b = set_read("Enter set b: ");
print_bits_hex("a = ", a);
print_bits_hex("b = ", b);
set_print("a = ", a);
set_print("b = ", b);
set_print("a union b = ", set_union(a, b));
set_print("a intersection b = ", set_intersection(a, b));
printf("cardinality(a) = %d\n", set_cardinality(a));
printf("is_member(42, a) = %d\n", (int)set_member(42, a));
```

source code for bitset.c

bitset.c: common set operations

```
set set_add(int x, set a) {  
    return a | ((set)1 << x);  
}  
  
set set_union(set a, set b) {  
    return a | b;  
}  
  
set set_intersection(set a, set b) {  
    return a & b;  
}  
  
set set_member(int x, set a) {  
    assert(x >= 0 && x < MAX_SET_MEMBER);  
    return a & ((set)1 << x);  
}
```

bitset.c: counting set members

```
int set_cardinality(set a) {  
    int n_members = 0;  
    while (a != 0) {  
        n_members += a & 1;  
        a >>= 1;  
    }  
    return n_members;  
}
```

```
set set_read(char *prompt) {  
    printf("%s", prompt);  
    set a = EMPTY_SET;  
    int x;  
    while (scanf("%d", &x) == 1 && x >= 0) {  
        a = set_add(x, a);  
    }  
    return a;  
}
```

bitset.c: set output

```
void set_print(char *description, set a) {
    printf("%s", description);
    printf("{");
    int n_printed = 0;
    for (int i = 0; i < MAX_SET_MEMBER; i++) {
        if (set_member(i, a)) {
            if (n_printed > 0) {
                printf(",");
            }
            printf("%d", i);
            n_printed++;
        }
    }
    printf("}\n");
}
```

Exercise: Bitwise Operations

Given the following variable declarations:

```
// a signed 8-bit value  
unsigned char x = 0x55;  
unsigned char y = 0xAA;
```

What is the value of each of the following expressions:

- $(x \ \& \ y)$ $(x \ \wedge \ y)$
- $(x \ \ll \ 1)$ $(y \ \ll \ 1)$
- $(x \ \gg \ 1)$ $(y \ \gg \ 1)$

Exercise: Bit-manipulation

Assuming 8-bit quantities and writing answers as 8-bit bit-strings:

What are the values of the following:

- 25, 65, ~ 0 , $\sim\sim 1$, 0xFF, $\sim 0xFF$
- $(01010101 \& 10101010)$, $(01010101 \mid 10101010)$
- $(x \& \sim x)$, $(x \mid \sim x)$

How can we achieve each of the following:

- ensure that the 3rd bit from the RHS is set to 1
- ensure that the 3rd bit from the RHS is set to 0