

Strategy for developing efficient programs:

1. Design the program well
 2. Implement the program well**
 3. Test the program well
 4. Only after you're sure it's working, *measure* performance
 5. If (and only if) performance is inadequate, *find* the "hot spots"
 6. *Tune* the code to fix these
 7. Repeat measure-analyse-tune cycle until performance ok
- (** see "Programming Pearls", "Practice of Programming", etc. etc.)

Rapid development of a prototype may be the best way to discover/assess performance issues.

Hence Fred Brooks maxim - "Plan To Throw One Away".

Where is execution time being spent?

Typically programs spend most of their execution time in a small part of their code.

This is often quoted as the 90/10 rule (or 80/20 rule or ...):

“90% of the execution time is spent in 10% of the code”

This means that

- most of the code has little impact on overall performance
- small parts of the code account for most execution time

We should clearly concentrate efforts at improving execution speed in the 10% of code which accounts for most of the execution time.

clang -p/gprof

Given the -p flag clang instruments a C program to collect profile information

When the program executes this data is left in the file gmon.out.

The program gprof analyzes this data and produces:

- number of times each function was called
- % of total execution time spent in the function
- average execution time per call to that function
- execution time for this function and its children

Arranged in order from most expensive function down.

It also gives a *call graph*, a list for each function:

- which functions called this function
- which functions were called by this function

Performance Improvement Example - Word Count

Program is slow on large inputs e.g.

```
$ clang -O3 word_frequency0.c -o word_frequency0
$ time word_frequency0 <WarAndPeace.txt >/dev/null
real    0m52.726s
user    0m52.643s
sys     0m0.020s
```

Performance Improvement Example - Word Count

We can instrument the program to collect profiling information and examine it with clang

```
$ clang -p -g word_frequency0.c -o word_frequency0_profile
$ head -10000 WarAndPeace.txt|word_frequency0_profile >/dev/null
$ gprof word_frequency0_profile
```

....

%	cumulative	self		self	total	
time	seconds	seconds	calls	ms/call	ms/call	name
88.90	0.79	0.79	88335	0.01	0.01	get
7.88	0.86	0.07	7531	0.01	0.01	put
2.25	0.88	0.02	80805	0.00	0.00	get_word
1.13	0.89	0.01	1	10.02	823.90	read_words
0.00	0.89	0.00	2	0.00	0.00	size
0.00	0.89	0.00	1	0.00	0.00	create_map
0.00	0.89	0.00	1	0.00	0.00	keys
0.00	0.89	0.00	1	0.00	0.00	sort_words

....

Performance Improvement Example - Word Count

Examine `{get}` and we find it traverses a linked list.

So replace it with a binary tree and the program runs 200x faster on *War and Peace*.

```
$ clang -O3 word_frequency1.c -o word_frequency1
$ time word_frequency1 <WarAndPeace.txt >/dev/null
real    0m0.277s
user    0m0.268s
sys     0m0.008s
```

Was C the best choice for our count words program?

Performance Improvement Example - Word Count

Shell, Perl and Python are slower - but a lot less code.

So faster to write, less bugs to find, easier to maintain/modify

```
$ time word_frequency1 <WarAndPeace.txt >/dev/null
```

```
real    0m0.277s
```

```
user    0m0.268s
```

```
sys     0m0.008s
```

```
$ time word_frequency.sh <WarAndPeace.txt >/dev/null
```

```
real    0m0.564s
```

```
user    0m0.584s
```

```
sys     0m0.036s
```

```
$ time word_frequency.pl <WarAndPeace.txt >/dev/null
```

```
real    0m0.643s
```

```
user    0m0.632s
```

```
sys     0m0.012s
```

```
$ time word_frequency.py <WarAndPeace.txt >/dev/null
```

```
real    0m1.046s
```

```
user    0m0.826s
```

Performance Improvement Example - cp - read/write

Here is a cp implementation in C using low-level calls to read/write

```
while (1) {
    char c[1];
    int bytes_read = read(in_fd, c, 1);
    if (bytes_read < 0) {
        perror("cp: ");
        exit(1);
    }
    if (bytes_read == 0)
        return;
    int bytes_written = write(out_fd, c, bytes_read);
    if (bytes_written <= 0) {
        perror("cp: ");
        exit(1);
    }
}
```


Performance Improvement Example - Fibonacci

Here is a simple Perl program to calculate the n-th Fibonacci number:

```
sub fib {  
    my ($n) = @_;  
    return 1 if $n < 3;  
    return fib($n-1) + fib($n-2);  
}  
printf "fib(%d) = %d\n", $_, fib($_) foreach @ARGV;
```

It becomes slow near n=35.

```
$ time fib0.pl 35  
fib(35) = 9227465  
real    0m10.776s  
user    0m10.729s  
sys     0m0.016s
```

we can rewrite in C.

Performance Improvement Example - Fibonacci

```
#include <stdio.h>

int fib(int n) {
    if (n < 3) return 1;
    return fib(n-1) + fib(n-2);
}

int main(int argc, char *argv[]) {
    for (int i = 1; i < argc; i++) {
        int n = atoi(argv[i]);
        printf("fib(%d) = %d\n", n, fib(n));
    }
}
```

Faster but the program's complexity doesn't change:

```
$ clang -O3 -o fib0 fib0.c
```

```
$ time fib0 45
```

```
fib(45) = 1134903170
```

```
real    0m4.994s
```

```
user    0m4.076s
```

Performance Improvement Example - Fibonacci

```
#!/usr/bin/perl -w
sub fib {
    my ($n) = @_;
    return 1 if $n < 3;
    $f{$n} = fib($n-1) + fib($n-2) if !defined $f{$n};
    return $f{$n};
}
printf "fib(%d) = %d\n", $_, fib($_) foreach @ARGV;
```

It is very easy to cache already computed results in a Perl hash.

This changes the program's complexity from exponential to linear.

```
$ time fib1.pl 45
fib(45) = 1134903170
real    0m0.004s
user    0m0.004s
sys     0m0.000s
```

Now for Fibonacci we could also easily change the program to an iterative