

# Balancing Search Trees

---

- Balancing Binary Search Trees
- Operations for Rebalancing
- Tree Rotation
- Insertion at Root
- Tree Partitioning
- Periodic Rebalancing
- Randomised BST Insertion
- An Application of BSTs: Sets

## ❖ Balancing Binary Search Trees

---

Observation: order of insertion into a tree affects its height

- **worst case**: keys inserted in ascending/descending order (effectively have a linked list, so search cost is  $O(n)$ )
- **best case** (for at-leaf insertion): keys inserted in pre-order (tree height  $\Rightarrow$  search cost is  $O(\log n)$ ; tree is balanced)
- **average case**: keys inserted in **random** order (tree height  $\Rightarrow$  search cost is  $O(\log n)$ ; but cost  $\geq$  best case)

Goal: build binary search trees which have

- minimum height  $\Rightarrow$  minimum worst case search cost

## ❖ ... Balancing Binary Search Trees

---

Perfectly-balanced tree with  $N$  nodes has

- $\forall$  nodes,  $\text{abs}(\text{\#nodes}(\text{LeftSubtree}) - \text{\#nodes}(\text{RightSubtree})) < 2$
- height of  $\log_2 N \Rightarrow$  worst case search  $O(\log N)$

Three *strategies* to improving worst case search in BSTs:

- **randomise** — reduce chance of worst-case scenario occurring
- **amortise** — do more work at insertion to make search faster
- **optimise** — implement all operations with performance bounds

## ❖ Operations for Rebalancing

---

To assist with rebalancing, we consider new operations:

Left rotation

- move right child to root; rearrange links to retain order

Right rotation

- move left child to root; rearrange links to retain order

Insertion at root


- each new item is added as the new root node

Partition

- rearrange tree around specified node (push it to root)

## ❖ Tree Rotation

Rotation operations:



Note: tree is ordered,  $t_1 < n_2 < t_2 < n_1 < t_3$

## ❖ ... Tree Rotation

---

Method for rotating tree T right:

- $n_1$  is current root;  $n_2$  is root of  $n_1$ 's left subtree
- $n_1$  gets new left subtree, which is  $n_2$ 's right subtree
- $n_1$  becomes root of  $n_2$ 's new right subtree
- $n_2$  becomes new root
- $n_2$ 's left subtree is unchanged

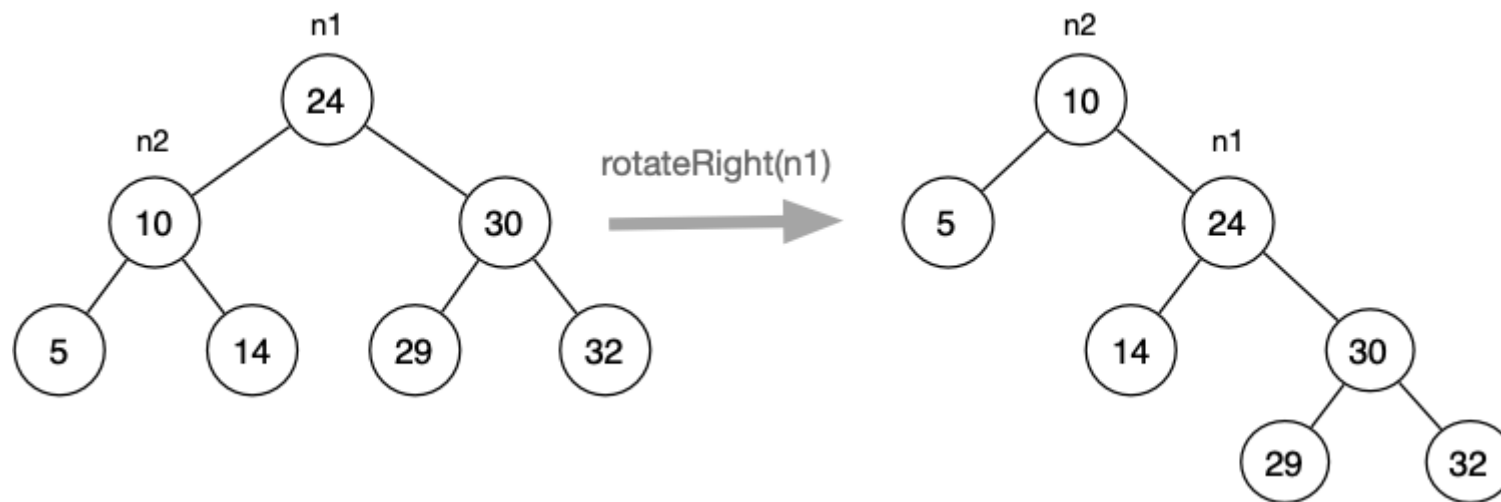
Left rotation: swap left/right in the above.

Rotation requires simple, localised pointer rearrangements

Cost of tree rotation:  $O(1)$

## ❖ ... Tree Rotation

Example of right rotation:



## ❖ ... Tree Rotation

---

Algorithm for right rotation:

```
rotateRight( $n_1$ ):  
|   Input   tree  $n_1$   
|   Output  $n_1$  rotated to the right  
|  
|   if  $n_1$  is empty  $\vee$  left( $n_1$ ) is empty then  
|       return  $n_1$   
|   end if  
|    $n_2 = \text{left}(n_1)$   
|   left( $n_1$ ) = right( $n_2$ )  
|   right( $n_2$ ) =  $n_1$   
|   return  $n_2$ 
```



## ❖ ... Tree Rotation

---

Algorithm for left rotation:

```
rotateLeft( $n_2$ ):  
|   Input   tree  $n_2$   
|   Output  $n_2$  rotated to the left  
|  
|   if  $n_2$  is empty  $\vee$  right( $n_2$ ) is empty then  
|       return  $n_2$   
|   end if  
|    $n_1 = \text{right}(n_2)$   
|   right( $n_2$ ) = left( $n_1$ )  
|   left( $n_1$ ) =  $n_2$   
|   return  $n_1$ 
```

## ❖ ... Tree Rotation

---

Cost considerations for tree rotation

- the rotation operation is cheap  $O(1)$
- if applied appropriately, will tend to improve tree balance

Sometimes rotation is applied from leaf to root, along one branch

- cost of this is  $O(\text{height})$
- payoff is improved balance which reduces height
- reduced height pushes search cost towards  $O(\log n)$

## ❖ Insertion at Root

---

Previous discussion of BSTs did insertion at leaves.

Different approach: insert new item at root.

Potential disadvantages:

- large-scale rearrangement of tree for each insert (apparently)

Potential advantages:

- recently-inserted items are close to root
- lower cost if recent items more likely to be searched

## ❖ ... Insertion at Root

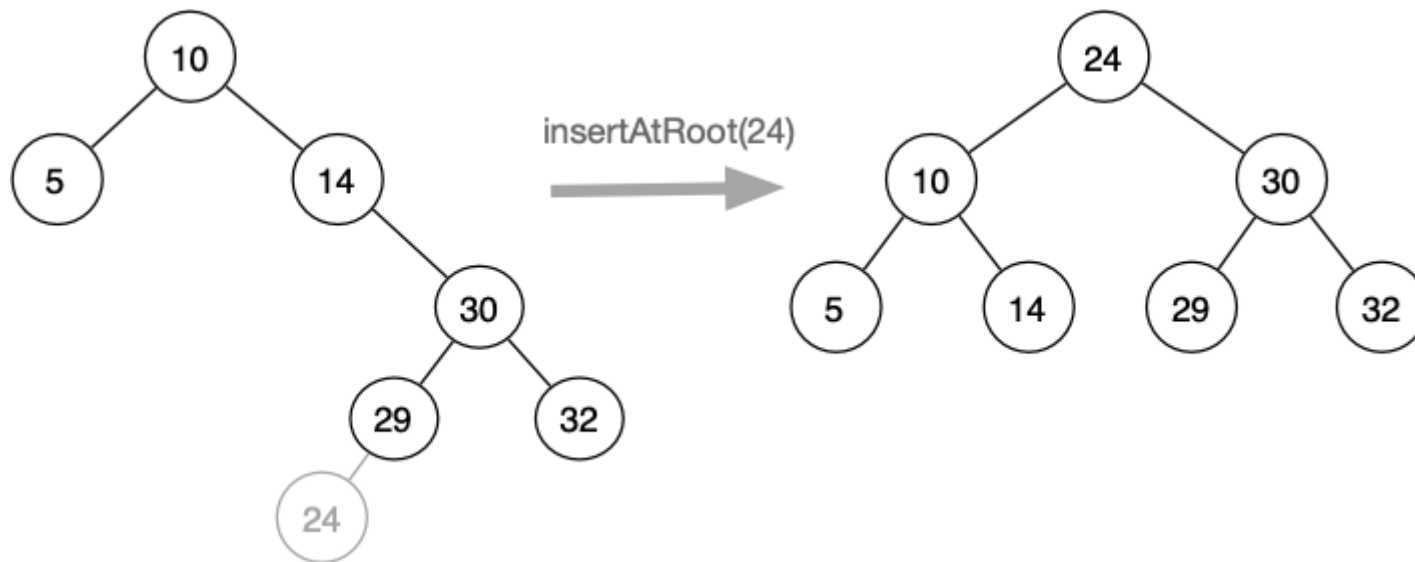
---

Method for inserting at root:

- base case:
  - tree is empty; make new node and make it root
- recursive case:
  - insert new node as root of appropriate subtree
  - lift new node to root by rotation

## ❖ ... Insertion at Root

Example of inserting at root:



## ❖ ... Insertion at Root

---

Algorithm for inserting at root:

```
insertAtRoot(t, it):  
|   Input tree t, item it to be inserted  
|   Output modified tree with item at root  
|  
|   if t is empty tree then  
|       t = new node containing item  
|   else if item < root(t) then  
|       left(t) = insertAtRoot(left(t), it)  
|       t = rotateRight(t)  
|   else if it > root(t) then  
|       right(t) = insertAtRoot(right(t), it)  
|       t = rotateLeft(t)  
|   end if  
|   return t;
```

## ❖ ... Insertion at Root

---


Analysis of insertion-at-root:

- same complexity as for insertion-at-leaf:  $O(\text{height})$ 
  - but cost is effectively doubled ... traverse down, rotate up
- tendency to be balanced, but no balance guarantee
- benefit comes in searching
  - for some applications, search favours recently-added items
  - insertion-at-root ensures these are close to root
- could even consider "move to root when found"
  - effectively provides "self-tuning" search tree

## ❖ Tree Partitioning

Tree partition operation **partition(tree, i)**

- re-arranges tree so that element with index  $i$  becomes root

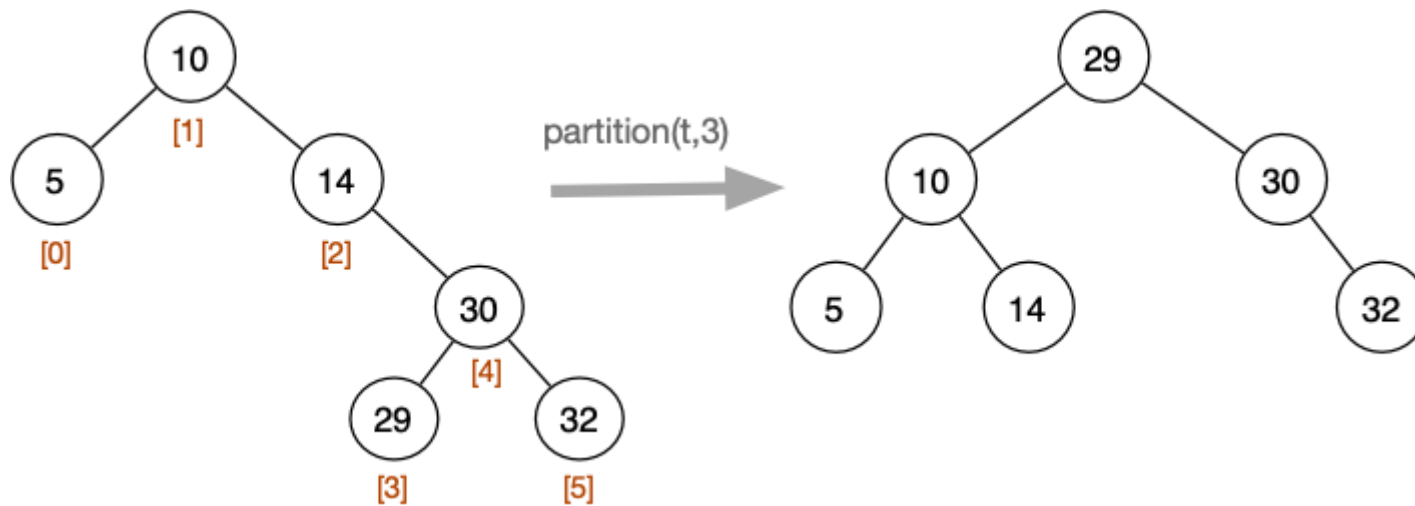


For tree with  $N$  nodes, indices are  $0..N-1$ , in LNR order



## ❖ ... Tree Partitioning

Example of partition:



## ❖ ... Tree Partitioning

---

Implementation of partition operation:

```
partition(tree,i):  
|   Input   tree with n nodes, index i  
|   Output tree with  $i^{\text{th}}$  item moved to the root  
|  
|   m=#nodes(left(tree))  
|   if i < m then  
|       left(tree)=partition(left(tree),i)  
|       tree=rotateRight(tree)  
|   else if i > m then  
|       right(tree)=partition(right(tree),i-m-1)  
|       tree=rotateLeft(tree)  
|   end if  
|   return tree
```

Note:  $\text{size}(\text{tree}) = n$ ,  $\text{size}(\text{left}(\text{tree})) = m$ ,  $\text{size}(\text{right}(\text{tree})) = n-m-1$

## ❖ ... Tree Partitioning

---

### Analysis of tree partitioning

- no requirement for search (using element index instead)
- after each recursive partitioning step, one rotation
- overall cost similar to insert-at-root

### Benefits

- tends to improve balance  $\Rightarrow$  improves search cost

## ❖ Periodic Rebalancing

---

An approach to maintaining balance:

- insert at leaves as before; periodically, rebalance the tree

```
|  Input  tree, item
|  Output tree with item randomly inserted
|
|  t=insertAtLeaf(tree,item)
|  if #nodes(t) mod k = 0 then
|      t=rebalance(t)
|  end if
|  return t
```

When to rebalance? e.g. after every  $k$  insertions

## ❖ ... Periodic Rebalancing

---

A problem with this approach ...


- operation `#nodes()` has to traverse whole (sub)tree
- to improve efficiency, change node structure

```
typedef struct Node {  
    int data;  
    int nnodes;          // #nodes in my tree  
    Tree left, right;    // subtrees  
} Node;
```

But maintaining `nnodes` requires extra work in other operations

## ❖ ... Periodic Rebalancing

How to rebalance a BST? Move median item to root.



## ❖ ... Periodic Rebalancing

---

Implementation of rebalance:

```
rebalance(t):  
|   Input   tree t with n nodes  
|   Output  t rebalanced  
|  
|   if  $n \geq 3$  then  
|   |   // put node with median key at root  
|   |   t=partition(t,  $\lfloor n/2 \rfloor$ )  
|   |   // then rebalance each subtree  
|   |   left(t)=rebalance(left(t))  
|   |   right(t)=rebalance(right(t))  
|   end if  
|   return t
```

## ❖ ... Periodic Rebalancing

---

Analysis of rebalancing: visits every node  $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every  $k$  insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely  $\Rightarrow$  Solution: real balanced trees (next week)



## ❖ Randomised BST Insertion

---

Reminder: order of insertion can dramatically affect shape of tree

Tree ADT has no control over order that keys are supplied.

We know that inserting in random order gives  $O(\log_2 n)$  search

Can the algorithm itself introduce some **randomness**?

In the hope that this randomness helps to balance the tree ...

## ❖ ... Randomised BST Insertion

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
|   Input   tree, item
|   Output tree with item randomly inserted
|
|   if tree is empty then
|       return new node containing item
|   end if
|   // p/q chance of doing root insert
|   if random() mod q < p then
|       return insertAtRoot(tree,item)
|   else
|       return insertAtLeaf(tree,item)
|   end if
```

E.g. 30% chance  $\Rightarrow$  choose  $p=3, q=10$

## ❖ ... Randomised BST Insertion

---

Cost analysis:

- similar to cost for inserting keys in random order:  $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
  - promote inorder successor from right subtree, OR
  - promote inorder predecessor from left subtree

## ❖ An Application of BSTs: Sets

---

Trees provide efficient search.

Sets require efficient search

- to find where to insert/delete
- to test for set membership

Logical to implement a set ADT via binary search tree.

## ❖ ... An Application of BSTs: Sets

---

Assuming we have BST implementation with type **Tree**

- which precludes duplicate key values
- which implements insertion, search, deletion

then **Set** implementation is


- **SetInsert(Set, Item)  $\equiv$  TreeInsert(Tree, Item)**
- **SetDelete(Set, Item)  $\equiv$  TreeDelete(Tree, Item.Key)**
- **SetMember(Set, Item)  $\equiv$  TreeSearch(Tree, Item.Key)**

What about union? and intersection?


## ❖ ... An Application of BSTs: Sets

---


Sets implemented via Trees:



After SetInsert(s,4):



After SetDelete(s,2):



## ❖ ... An Application of BSTs: Sets

---

Concrete representation:

```
#include <Tree.h>

typedef struct SetRep {
    int    nelems;
    Tree   root;
} SetRep;

Set newSet() {
    Set S = malloc(sizeof(SetRep));
    assert(S != NULL);
    S->nelems = 0;
    S->root = newTree();
    return S;
}
```



