

Relational Algebra

- Relational Algebra
- Notation
- Describing RA Operations
- Example Database #1
- Example Database #2
- Rename
- Selection
- Projection

◆ Relational Algebra

Relational algebra (RA) can be viewed as ...

- mathematical system for manipulating relations, or
- data manipulation language (DML) for the relational model

Relational algebra consists of:

- **operands**: relations, or variables representing relations
- **operators** that map relations to relations
- rules for combining operands/operators into expressions
- rules for evaluating such expressions

Why is it important?

- because it forms the basis for DBMS implementation
- relational algebra ops are like the machine code for DBMSs

◆ Relational Algebra (cont)

Core relational algebra operations:


- **rename**: change names of relations/attributes
- **selection**: choosing a subset of tuples/rows
- **projection**: choosing a subset of attributes/columns
- **union, intersection, difference**: combining relations
- **product, join**: combining relations

Common extensions include:

- **aggregation, projection++, division**

❖ Relational Algebra (cont)

Select, project, join provide a powerful set of operations for building relations and extracting interesting data from them.



Adding set operations and renaming makes RA **complete**.

❖ Notation

Standard treatments of relational algebra use Greek symbols.

We use the following notation (because it is easier to reproduce):

Operation	Standard Notation	Our Notation
Selection	$\sigma_{expr}(Rel)$	$Sel[expr](Rel)$
Projection	$\pi_{A,B,C}(Rel)$	$Proj[A,B,C](Rel)$
Join	$Rel_1 \bowtie_{expr} Rel_2$	$Rel_1 \text{ Join}[expr] Rel_2$
Rename	$\rho_{schema} Rel$	$Rename[schema](Rel)$

For other operations (e.g. set operations) we adopt the standard notation.
Except when typing in a text file, where * = intersection, + = union

❖ Describing RA Operations

We define the semantics of RA operations using

- "conditional set" expressions e.g. $\{X / \text{condition on } X\}$
- tuple notations:
 - $t[AB]$ (extracts attributes A and B from tuple t)
 - (x,y,z) (enumerated tuples; specify attribute values)
- quantifiers, set operations, boolean operators

For each operation, we also describe it operationally:

- give an algorithm to compute the result, tuple-by-tuple
- the algorithm is not generally how it will be computed in practice

❖ Describing RA Operations (cont)

All RA operators return a result of type **relation**.

For convenience, we can name a result and use it later.

E.g.

```
Temp = R op1 S op2 T
Res  = Temp op3 Z
-- which is equivalent to
Res  = (R op1 S op2 T) op3 Z
```

Each "intermediate result" has a well-defined schema.

❖ Example Database #1

R

A	B	C	D
a	1	x	4
b	2	y	5
c	4	z	4
d	8	x	5
e	1	y	4
f	2	x	5

S

D	E	F
1	a	x
2	b	y
3	c	x
4	a	y
5	b	x

❖ Example Database #2

Beers(name,manf) (VB, Carlton) (New, Tooheys) (Porter, Maltshovel) ...	Beers(name,addr,licence) (CBH, Coogee,433122) (Royal, Randwick, 632987) (Regent, Kingsford,112112) ...	Frequents(drinker,bar) (John, CBH) (Gernot, CBH) (Gernot, Regent) ...
Likes(drinker, beer) (Andrew, New) (Gernot, Porter) (John, Pale Ale) ...	Beers(name,addr,phone) (John, Alexandria, 93111139) (Gernot, Newtown, 92422429) (Andrew, Glebe, 90411049) ...	Sells(beer,bar,price) (CBH, New, 2.50) (CBH, VB, 1.99) Royal, Porter, 3.00 ...

◆ Rename

Rename provides "schema mapping".

If expression E returns a relation $R(A_1, A_2, \dots, A_n)$, then

$\text{Rename}[S(B_1, B_2, \dots, B_n)](E)$

gives a relation called S

- containing the same set of tuples as E
- but with the name of each attribute changed from A_i to B_i

Rename is like the identity function on the *contents* of a relation

The only thing that Rename changes is the schema.

❖ Rename (cont)

Rename can be viewed as a "technical" apparatus of RA.

We can also use implicit rename/project in sequences of RA operations, e.g.

```
-- R(a,b,c), S(c,d)
Res = Rename[Res(b,c,d)](Project[b,c](Sel[a>5](R)) Join S)
-- VS
Tmp1 = Select[a>5](R)
Tmp2 = Project[b,c](Tmp1)
Tmp3 = Rename[Tmp3(cc,d)](S)
Tmp4 = Tmp2 Join[c=cc] Tmp3
Res = Rename[Res(b,c,d)](Tmp4)
-- VS
Tmp1(b,c) = Select[a>5](R)
Tmp2(cc,d) = S
Res(b,c,d) = Tmp1 Join[c=cc] Tmp2
```

In SQL, we achieve a similar effect by defining a set of views

◆ Selection

Selection returns a subset of the tuples in a relation $r(R)$ that satisfy a specified condition C .

$$\sigma_C(r) = \text{Sel}[C](r) = \{ t \mid t \in r \wedge C(t) \}$$

C is a boolean expression on attributes in R .

Result size: $|\sigma_C(r)| \leq |r|$

Result schema: same as the schema of r (i.e. R)

Algorithmic view:

```
result = {}  
for each tuple  $t$  in relation  $r$   
    if ( $C(t)$ ) { result = result  $\cup$  { $t$ } }
```

❖ Selection (cont)

Examples of selection:

R

A	B	C	D
a	1	x	4
b	2	y	5
c	4	z	4
d	8	x	5
e	1	y	4
f	2	x	5

$\text{Sel}[A=C](R)$

A	B	C	D
---	---	---	---

$\text{Sel}[B=1](R)$

A	B	C	D
a	1	x	4
e	1	y	4

$\text{Sel}[B \geq D](R)$

A	B	C	D
c	4	z	4
d	8	x	5

$\text{Sel}[A=b \text{ or } A=c](R)$

A	B	C	D
b	2	y	5
c	4	z	4

◆ Selection (cont)

Querying with relational algebra (selection) ...

- Details of all bars in The Rocks

```
Result = Sel[addr=The Rocks](Bars)
```

- Beers made by Sierra Nevada

```
SNBeers = Sel[manf=Sierra Nevada](Beers)  
Result  = Rename[beer](Proj[name](SNBeers))
```

◆ Projection

Projection returns a set of tuples containing a subset of the attributes in the original relation.

$$\pi_X(r) = \text{Proj}[X](r) = \{ t[X] \mid t \in r \}, \text{ where } r(R)$$

X specifies a subset of the attributes of R .

Note that removing key attributes can produce duplicates.

In RA, duplicates are removed from the result **set**.

(In RDBMS's, duplicates are retained (i.e. they use bags, not sets))

Result size: $|\pi_X(r)| \leq |r|$ Result schema: $R'(X)$

Algorithmic view:

```
result = {}  
for each tuple t in relation r  
    result = result ∪ {t[X]}
```


◆ Projection (cont)

Examples of projection:

R

A	B	C	D
a	1	x	4
b	2	y	5
c	4	z	4
d	8	x	5
e	1	y	4
f	2	x	5

Proj[A,B,C](R)

A	B	C
a	1	x
b	2	y
c	4	z
d	8	x
e	1	y
f	2	x

Proj[B,D](R)

B	D
1	4
2	5
4	4
8	5

Proj[D](R)

D
4
5

◆ Projection (cont)

Querying with relational algebra (projection)...

- Names of all beers

```
Result = Proj[name](Beers)
```

- Names of drinkers who live in Newtown

```
Result = Proj[name](Sel[addr=Newtown](Drinkers))
```

- What are all of the breweries?

```
Result(brewer) = Proj[manf](Beers)
```

Produced: 6 Nov 2020